

Summary

Johnson ODEJIDE

2023-03-05

Contents


Plot the estimated regression function	1
Fit a regression line	2
Estimate a slope	3
Find specific residual for $X_i = 3$. Show all of your work . Determine if the observed value of 22 is above or below average.	3
Verify that your fitted regression line goes through the point (\bar{X}, \bar{Y})	3
Interpretation of the Multiple R-Squared	4
Computation of confidence interval in R	4
Raw computation of confidence interval	4
Difference between confidence interval and prediction interval	5
Prediction Interval in R	5
Confidence Interval in Prediction	5
Plotting Residuals using QQplot	5
Residual plots (Scatter plot)	6
Residual against Fitted Plot	7
To reject null hypotheses or not?	8
Boxcox	9
SSE from lambda results	11
Initializing Matrices	13
Transpose of Matrices	14
Inverse of Matrices	14
Multiplication of uneven dimensions	14
Getting the dimensions of Matrices	14
Intercept and Slope using Matrices	14
Fitted values using Matrices	15
Residuals using Matrices	15

Plot the estimated regression function

```
route <- c(1, 0, 2, 0, 3, 1, 0, 1, 2, 0)
ampules <- c(16, 9, 17, 12, 22, 13, 8, 15, 19, 11)

df = tibble(route, ampules)

df %>%
  ggplot(aes(x = route, y = ampules)) +
  geom_point(color = "red") +
  geom_smooth(method = "lm", se = F, formula = y ~ x)
```



Fit a regression line

```
lm.fit <- lm(ampules ~ route)
summary(lm.fit)
```

```
##
## Call:
## lm(formula = ampules ~ route)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
##     -2.2     -1.2      0.3      0.8      1.8
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  10.2000    0.6633   15.377 3.18e-07 ***
## route         4.0000    0.4690    8.528 2.75e-05 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.483 on 8 degrees of freedom
## Multiple R-squared:  0.9009, Adjusted R-squared:  0.8885
## F-statistic: 72.73 on 1 and 8 DF, p-value: 2.749e-05
```

Estimate a slope

Estimate the increase in the expected number of ampules broken when there are 2 transfers as compared to 1 transfer. (Hint: Use the slope of the model)

```
# The formula is given as
```

```
#  $\hat{y} = 10.2 + 4x$ 
```

```
# when  $x = 1$ 
```

```
yhat_1 <- 10.2 + 4*1
```

```
yhat_1
```

```
## [1] 14.2
```

```
# Calculate for when  $X = 2$ 
```

```
yhat_2 <- 10.2 + 4*2
```

```
# Estimate the increase
```

```
slope <- yhat_2 - yhat_1
```

```
slope
```

```
## [1] 4
```

Find specific residual for $X_i = 3$. Show all of your work . Determine if the observed value of 22 is above or below average.

```
# For  $x = 3$ 
```

```
#  $\hat{y} = 10.2 + 4*3$ 
```

```
#  $\hat{y} = 10.2 + 12 = 22.2$ 
```

```
# residual for  $X_3 = 22.2 - 22 = 0.2$  (Below average because it is negative)
```

```
yhat <- 10.2 + 4 * 3
```

```
residual <- 22 - yhat # Where 22 is the observed value
```

```
residual
```

```
## [1] -0.2
```

Conclusion: Since the residual is negative, we conclude that it is below average

Verify that your fitted regression line goes through the point (\bar{X} , \bar{Y})

```
Xbar = mean(route)
```

```
Ybar = mean(ampules)
```

```
# To verify that the fitted regression line goes through the point, we substitute  $x$  in the equation for
```

```
yhat <- 10.2 + 4 * 1
```

```
print(paste("yhat = ", yhat))
```

```
## [1] "yhat = 14.2"
```

```
print(paste("Xbar = ", Xbar, ", Ybar = ", Ybar, "Yhat at points(1, 14.2) = ", yhat))
```

```
## [1] "Xbar = 1 , Ybar = 14.2 Yhat at points(1, 14.2) = 14.2"
```

Remark:

The fitted regression line goes through the point at $\bar{x} = 1$ and $\bar{y} = 14.2$ since \hat{y} and \bar{y} are the same value.

Interpretation of the Multiple R-Squared

Which value in the R summary output table determines if your model is doing a good job explaining the variation in the dependent variable produced by the model. In this case identify this specific proportion of variation.

The Multiple R-squared does the job of explaining the variation in the dependent variable produced by the model.

In this case, the Multiple R-Squared is 0.9009 meaning that about 90% of the variability in the dependent variable (ampules) can be explained by the model.

Computation of confidence interval in R

```
X = c(1, 0, 2, 0, 3, 1, 0, 1, 2, 0)
Y = c(16, 9, 17, 12, 22, 13, 8, 15, 19, 11)

d.data <- tibble(X, Y)

conf.95 <- qt(p=.025, df=8, lower.tail = FALSE)

# conf.95

lm.fit <- lm(Y ~ X, data = d.data)
# lm.fit

# summary(lm.fit)

# qt(p=.025, df=8, lower.tail = FALSE)
# 4.00 +/- 2.306(0.469)

upper_bound <- 4.00 + 2.306 * 0.469
lower_bound <- 4.00 - 2.306 * 0.469

conf.int <- c(lower_bound, upper_bound)
conf.int

## [1] 2.918486 5.081514
```

Raw computation of confidence interval

```
se = 0.4690
df = 8
b +/- t(se)
t = 2.306004
4.0 +/- 2.306004(0.4690)
confidence interval = (4.0 - 2.306004(0.4690), 4.0 + 2.306004(0.4690))
= (2.9185, 5.0815)
```

Difference between confidence interval and prediction interval

While the prediction interval predicts in what range a future observation will fall, the confidence interval shows in what range of values the prediction falls based on some data provided already. In summary, confidence interval predicts what is available within the limits of the data while prediction interval is able to predict the future.

Prediction Interval in R

```
new_df <- data.frame(X = 19)

predict(object = lm.fit, newdata = new_df, interval = "prediction") %>%
cbind(new_df)

##      fit      lwr      upr  X
## 1 86.2 66.40325 105.9967 19
```

Confidence Interval in Prediction

```
predict(object = lm.fit, newdata = new_df, interval = "confidence") %>%
cbind(new_df)

##      fit      lwr      upr  X
## 1 86.2 66.70097 105.699 19
```


Remarks on Prediction and Confidence Interval

Prediction Interval when $X = 19$ (66.4033, 105.9967)

Confidence Interval when $X = 19$ (66.701, 105.699)

Plotting Residuals using QQplot


```
qqnorm(lm.fit$residuals)
qqline(lm.fit$residuals)
```



Residual plots (Scatter plot)

```
lm.resid <- resid(lm.fit)


plot(d.data$X , lm.resid,
     xlab = "X",
     ylab = "residuals",
     main = "Residual plot")
abline(0, 0)
```



Residual against Fitted Plot

```
lm.fitted <- fitted(lm.fit)

plot(lm.fitted, lm.resid)
abline(0, 0)
```



To reject null hypotheses or not?

```
x <- c(9,9,9,7,7,7,5,5,5,3,3,3,1,1,1)
y<- c(.07,.09,.08,.16,.17,.21,.49,.58,.53,1.22,1.15,1.07,2.84,2.57,3.10)
```

```
linear <- lm(y ~ x)
summary(linear)
```

```
##
## Call:
## lm(formula = y ~ x)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.5333 -0.4043 -0.1373  0.4157  0.8487
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   2.5753     0.2487  10.354 1.20e-07 ***
## x             -0.3240     0.0433  -7.483 4.61e-06 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.4743 on 13 degrees of freedom
## Multiple R-squared:  0.8116, Adjusted R-squared:  0.7971
## F-statistic: 55.99 on 1 and 13 DF,  p-value: 4.611e-06
```



```

anova(linear)

## Analysis of Variance Table
##
## Response: y
##           Df Sum Sq Mean Sq F value    Pr(>F)
## x           1 12.5971   12.597   55.994 4.611e-06 ***
## Residuals  13  2.9247    0.225
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

f.value <- 55.994

f.critical <- qf(p = 0.025, df1 = 1, df2 = 13, lower.tail = FALSE)

print(paste("F value = ", f.value, "F Critical = ", f.critical))

## [1] "F value = 55.994 F Critical = 6.41425430025058"
ifelse(f.value > f.critical, "Reject the null hypothesis", "Fail to reject the null hypothesis")

## [1] "Reject the null hypothesis"

```

Boxcox

```


library(MASS)

##
## Attaching package: 'MASS'

## The following object is masked from 'package:dplyr':
##
##      select

box_cox <- boxcox(Y ~ X)

```



```
# Optimum lambda
```

```
lambda <- box_cox$x[which.max(box_cox$y)]
lambda
```

```
## [1] 1.717172
```

```
new_model <- lm(((Y^lambda-1)/lambda) ~ X)
new_model
```

```
##
```

```
## Call:
```

```
## lm(formula = ((Y^lambda - 1)/lambda) ~ X)
```

```
##
```

```
## Coefficients:
```

```
## (Intercept)          X
```

```
##      29.99      27.88
```

```
lambda.3 <- lm(((Y^.3 - 1) / .3) ~ X)
```

```
anova(lambda.3)
```

```
## Analysis of Variance Table
```

```
##
```

```
## Response: ((Y^0.3 - 1)/0.3)
```

```
##      Df Sum Sq Mean Sq F value    Pr(>F)
```

```
## X      1  3.8197   3.8197  47.646 0.0001242 ***
```

```
## Residuals  8  0.6413   0.0802
```

```
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
lambda.5 <- lm(((Y^.5 - 1) / .5) ~ X)
anova(lambda.5)
```

```
## Analysis of Variance Table
##
## Response: ((Y^0.5 - 1)/0.5)
##           Df Sum Sq Mean Sq F value    Pr(>F)
## X           1 11.0405  11.0405   54.078 7.965e-05 ***
## Residuals   8  1.6333   0.2042
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Remark: An Appropriate transformation of Y would be $Y^{(0.5050505)}$

SSE from lambda results

SSE = 0.584 when lamda = .3 suggested transformation is $Y^{0.3}$

SSE = 1.51 when lamda = .4 suggested transformation is $Y^{0.4}$


SSE = 4.20 when lamda = .5 suggested transformation is $Y^{0.5}$

```
Y_trans <- lm(Y^0.5 ~ X)
summary(Y_trans)
```

```
##
## Call:
## lm(formula = Y^0.5 ~ X)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.37222 -0.12632  0.01059  0.13922  0.27399
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  3.20064    0.10103   31.679 1.07e-09 ***
## X            0.52537    0.07144    7.354 7.96e-05 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.2259 on 8 degrees of freedom
## Multiple R-squared:  0.8711, Adjusted R-squared:  0.855
## F-statistic: 54.08 on 1 and 8 DF,  p-value: 7.965e-05
```


Remark: $\text{sqrt}(\hat{y}) <- 10.26093 + 1.07629X$

```
plot(X, sqrt(Y))
abline(lm(sqrt(Y)~X))
```



Residual Plot

```
plot(fitted(Y_trans), residuals(Y_trans))  
abline(0,0, col = "red")
```



Initializing Matrices

```
A <- matrix(c(1, 4,
              2, 6,
              3, 8), ncol = 2, byrow = TRUE)

B <- matrix(c(1, 3,
              1, 4,
              2, 5), ncol = 2, byrow = TRUE)

C <- matrix(c(3, 8, 1,
              5, 4, 0), ncol = 3, byrow = TRUE)

D <- matrix(c(5, 3,
              15, 6), ncol = 2, byrow = TRUE)

D
```

```
##      [,1] [,2]
## [1,]    5    3
## [2,]   15    6
```

Transpose of Matrices

```
t(B)
```

```
##      [,1] [,2] [,3]
## [1,]    1    1    2
## [2,]    3    4    5
```

Inverse of Matrices

Remark: To inverse a matrix, it has to square!

```
solve(D)
```

```
##      [,1]      [,2]
## [1,] -0.4  0.2000000
## [2,]  1.0 -0.3333333
```

Multiplication of uneven dimensions

```
A %*% C -> mult.res
mult.res
```

```
##      [,1] [,2] [,3]
## [1,]   23   24    1
## [2,]   36   40    2
## [3,]   49   56    3
```

Getting the dimensions of Matrices

```
dim(mult.res)
```

```
## [1] 3 3
```

Intercept and Slope using Matrices

```
Y <- matrix(c(124,
              95,
              71,
              45,
              18), ncol = 1, byrow = TRUE)

X <- matrix(c(1, 49,
              1, 69,
              1, 89,
              1, 99,
              1, 109), ncol = 2, byrow = TRUE)

t(X) -> transposeX
# transposeX

dim(X)

## [1] 5 2
```

```
transposeX%*%X -> Product2
# Product2
```

```
det(Product2)
```

```
## [1] 11600
```

```
solve(Product2)
```

```
##           [,1]      [,2]
## [1,]  3.16939655 -0.0357758621
## [2,] -0.03577586  0.0004310345
```

```
interceptandslope <- solve(Product2)%*%transposeX%*%Y
```

```
interceptandslope
```

```
##           [,1]
## [1,] 211.270690
## [2,] -1.694828
```

Remark: Intercept = 211.27, and Slope = -1.6948

Fitted values using Matrices

```
X %*% interceptandslope
```

```
##           [,1]
## [1,] 128.22414
## [2,]  94.32759
## [3,]  60.43103
## [4,]  43.48276
## [5,]  26.53448
```

Residuals using Matrices

```
Y - X %*% interceptandslope
```

```
##           [,1]
## [1,] -4.2241379
## [2,]  0.6724138
## [3,] 10.5689655
## [4,]  1.5172414
## [5,] -8.5344828
```